


Lecture 1

- Lecture 1: https://www.youtube.com/watch?v=v6w_zVHL8WQ
- Written by Milad Abdi with some AI assistance.
- These notes are meant as a reference and cannot replace watching the lecture.
- **Skill Stack**
 - Base: Perception, planning, and control
 - **Machine Learning layer** (vision and perception algorithms): Used when physics-based models are insufficient (e.g., drifting)
- Formal Methods: To prove safety mathematically
- F1-Tenth uses NVIDIA Jetson Xavier NX is a GPGPUs (General-Purpose computing on Graphics Processing Units):
 - Supports running and programming applications **directly on the device**.
 - Can connect a **monitor, keyboard, and mouse** (no external laptop required).
 - Can also be accessed **remotely via SSH over Ethernet**.
 - Enabled Workflow
 - Develop locally (Mac/Windows, Docker or not)
 - Push to GitHub
 - SSH into Jetson
 - Pull repo
 - Build/run **in the Jetson's Linux environment**
- **Lesson Plan**
 - Learn to drive: Avoid Crashing (Reactive planning)
 - Learn to race: Follow the Raceline (Purse pursuit), Race & Overtake (Graph planner)
- **Balancing Performance and Safety:** Current AV (autonomous vehicles) technology still **struggles in non-cooperative scenarios** like merging due to **competing objectives**:
 - Max performance: Negotiate the merge without delay or hesitation
 - Max safety: Avoid catastrophic failures and crashes
 - **Racing (autonomously) highlights this performance safety tradeoff. This is why we race.**
 1. **Detecting vehicle limits:** Unstructured environment; different tracks and conditions; different vehicle setups.
 2. **Decision making at the vehicle limits:** High prediction uncertainty; strategy, energy, overtaking.
 3. **Handling at the vehicle limits:** High speeds and acceleration; high planning horizon necessary; small reaction times

- If we only trained in civilian driving then we would be too safe and not push the limits of AV technology. The civilian cars trickle down from the race cars.
- **AV Perception, Planning, Control (PPC) pipeline:**
 - Sensors & Perception: To begin navigating
 - Camera feeds (RGB images: color + texture). Texture is beyond color e.g., asphalt is rough and grainy but walls are smooth or brick.
 - Lidar feeds (3D point clouds: precise distance measurements)
 - Capture depth views (per-pixel distance from camera/LiDAR)
 - Capture optical flow views (relative motion of pixels between frames for **speed & direction**)
 - Object detection (what objects are present and where they are)
 - Object tracking (maintaining object identity and motion over time)
 - Semantic Segmentation (pixel-level classification of scene elements). Make every pixel assigned a class label e.g., this pixel is road, this pixel is car, this pixel is pedestrian.
 - Hierarchical Planner: Different levels of planner that repeats 20-100 times a second.
 - Note that each planner repeats at different speeds. Mission planner doesn't repeat nearly as much as local planner.
 1. Mission Planner: Main goal. E.g., Go from Science World to SFU, what streets to go on.
 2. Behavioural Planner: Choices at each step. E.g., Drive or Pause
 3. Local Planner: Act. E.g., this is the best trajectory -> generate sequence of steering and acceleration inputs to follow that path



Everything combined:
a function that *generates* a sequence of ***steering*** and ***acceleration*** inputs...

- **Pipeline Continuation:** Hardware (sensors) -> Software (PPC) -> Hardware (race car)
- **Perception:**
 - Perception-only based navigation: No bigger plan for where to go, just driving based off reactions (just reactions to sensor input).
 - Localization: Estimating the vehicle's pose (position + orientation) in the world
 - Detection: Identifying objects and obstacles in the environment
- **Planning:**
 - Prediction: Estimating future motion / path
 - Behavior: High-level decisions (e.g., how to handle other drivers or sequence turns)
 - Trajectory: Generating the right trajectories that minimize cost while adhering to dynamic systems constraints.
- **Control:**
 - Path & Velocity: Computing steering and speed commands to follow a planned path
 - PID control: Feedback controller using proportional, integral, and derivative terms to reduce error
 - Modern predictive control: Optimization-based control (e.g., MPC) that accounts for future states
- **Types of Failures:**

- False Positive
- False Negative
- **F1Tenth Gym:**
 - Lightweight 2D simulator built in Python
 - Asynchronous
 - Faster than real time execution (30x realtime)
 - Realistic vehicle simulation and collision
 - Publishes laser scan and odometry data
 - Built for fast prototyping
- **LG SVL Simulator:**
 - 3D simulator with unity based environments
 - **Distributed cloud simulation**
 - Bridges: ROS, ROS2, Apollo, Autoware, Python
 - Sophisticated physics simulation
 - Sensor integration
- A **bridge** connects **different systems or software**, allowing them to communicate seamlessly.
 - Connects the simulator (e.g., LG SVL) to nodes (e.g., from ROS2/Autoware).
 - Translates **sensor data** (LiDAR, camera, GPS, IMU) from the simulator into **ROS2 messages**.
 - Sends **control commands** from the software back to the simulator.

Future Module Insights

- **Race 1: Reactive Methods**
 - **Pose Representation and Transforms:**
 - Each sensor provides measurements in the frame of reference specific to that sensor.
 - **Sensor Fusion Algorithms:** Combine data from multiple sensors to get a **more accurate and reliable understanding of the environment** than any single sensor alone.
 - Outcome: Coordinate frames, Rigid body transforms.
 - **Multiple Reference Frames:** World makes sense if we put laser scans in the **global map frame** instead of the **laser frame** -> Gives more useful info when planning in **global frame** (a consistent world-fixed coordinate system) than in the **observation frame** (the frame where each individual sensor observation is interpreted).
 - **Electronic Speed Control:** Actuate vehicle via commands in physics units.

- Outcome: PID (Proportional-Integral-Derivative controller for precise speed/position control), motor control (low-level signals sent to the motor to achieve desired motion), etc.
- **Walk Following:** How we can drive the car around the track.
 - Outcome: Basics of PID, how to compute error, and failure modes.
- **Obstacle performance:** How can we avoid obstacles.
 - Outcome: Basics of reactive navigation, avoidance on both static and dynamic obstacles.
- **Race 2: Map-based Methods**
 - **Scan Matching:** Where is the robot with respect to the previous frame.
 - Outcome: Iterative closest point algorithm, implementing a real research paper.
 - Scan matching is a fundamental localization algorithm, and is used in most of the modern SLAM algorithms.
 - **SLAM (simultaneous localization and mapping):** Method that lets you build a map and localize your vehicle in that map at the same time. Used for path planning and obstacle avoidance. Learns the map -> Exploits the learned map for speed
 - **Simultaneous Localization and Mapping with Cartographer:** How to use state-of-the-art tools for map building
 - Outcome: Understand the Cartographer paper and how it relates to scan matching.
 - **Localization:** Estimate the robot's pose within a map using sensor data, often by matching LiDAR scans or other observations to the map.
 - **Localization: Particle Filter:** Given a map of the world and multiple sensor observations, what is the pose of my robot?
 - Outcome: Understanding particle filter, which is a version of a bayesian filter.
 - **Pure Pursuit:** How to track a reference trajectory given a map and the ability to localize?
 - Outcome: Closed form geometric approach and alternatives.
- **Race 3: Head-to-Head**
 - **Motion Planning:** How do we combine the capabilities of map based methods while being able to avoid obstacles
 - Outcome: Understanding search-based motion planning, probabilistic planning methods, RRT and its variants.
 - **Occupancy Grid:** Approximating the real world with a discrete representation, also relates back to SLAM lecture
 - **Planning in discrete space:** With search-based planning methods (A*, Dijkstra's)

- **Planning in continuous space:** With probabilistic planning methods (RRT, RRT*)
- **Model Predictive Control:** Create dynamically feasible trajectories for overtaking.
 - Outcome: Trajectory optimization & Sampling based MPC
- **Learning-based CV (Computer Vision): Detection and Pose Estimation:** Where is the other car without using fiducial markers?
 - Outcome: Understanding the multi-view geometry, the epipolar constraint, stereo vision, and using Convolution Neural Network detectors.
- **Classical CV: Detection and Pose Estimation:** Where is the other car?
 - Outcome: Understanding camera model, single view geometry, Homography, detecting features, and prediction.
- **RL (Reinforcement Learning):** How to learn from human drivers (RL can be more than this but this course is focused on human imitation).
 - Outcome: Understand imitation learning and implement it.