

Tutorial 1

- Tutorial 1: <https://www.youtube.com/watch?v=EU-QaO6xTv4>
- Written by Milad Abdi

Docker

- Allows for easy deployment of software in a loosely isolated sandbox (containers).
- Can run many containers simultaneously.
- Containers are **lightweight** and contain everything needed to run the application.
 - This fixes "but it worked on my machine" problems.
- Allows us to package an application with **all its dependencies into a standardized unit**.
- **Architecture:**
 - Client: You can use docker on the command line (terminal) or use Docker Desktop.
 - Serverside: Handles all containers and images.
 - Registry (DockerHub): Place for publicly available containers and images.
- **Images:**
 - Read-only template with instructions for creating a Docker container. (Instructions to create your sandbox with whatever software you want installed).
 - Image often based on another image. E.g., Your ROS 2 image would have an image of Ubuntu with ROS 2 added onto it.
 - You can create your own images or use images made by others.
- **Building Images:**
 - Create `Dockerfile` and write the instructions syntax which define the steps needed to **create the image** and **run it**.
 - **Layer caching:** Each instruction creates a layer in the image. If you change your Dockerfile, only the changed layers will be rebuilt. Thus, Docker is lightweight and fast.
- **Containers:**
 - **Runnable instance of an image** (you can have multiple instances of the same container at once).
 - You can create, start, stop, move, or delete a container.
 - You can connect a container to one or more networks (to let them talk to each other or talk to you separately), and attach storage to it.
 - Container is **well isolated** from other containers and host machine. You can control how much isolation between your container and other containers or host machine.
 - Container is **defined by its image** and the **configuration options** you give it when you **create or start** the container.

- Container removed -> Any changes to its state that are **not** stored in persistent storage disappear.

Docker Commands

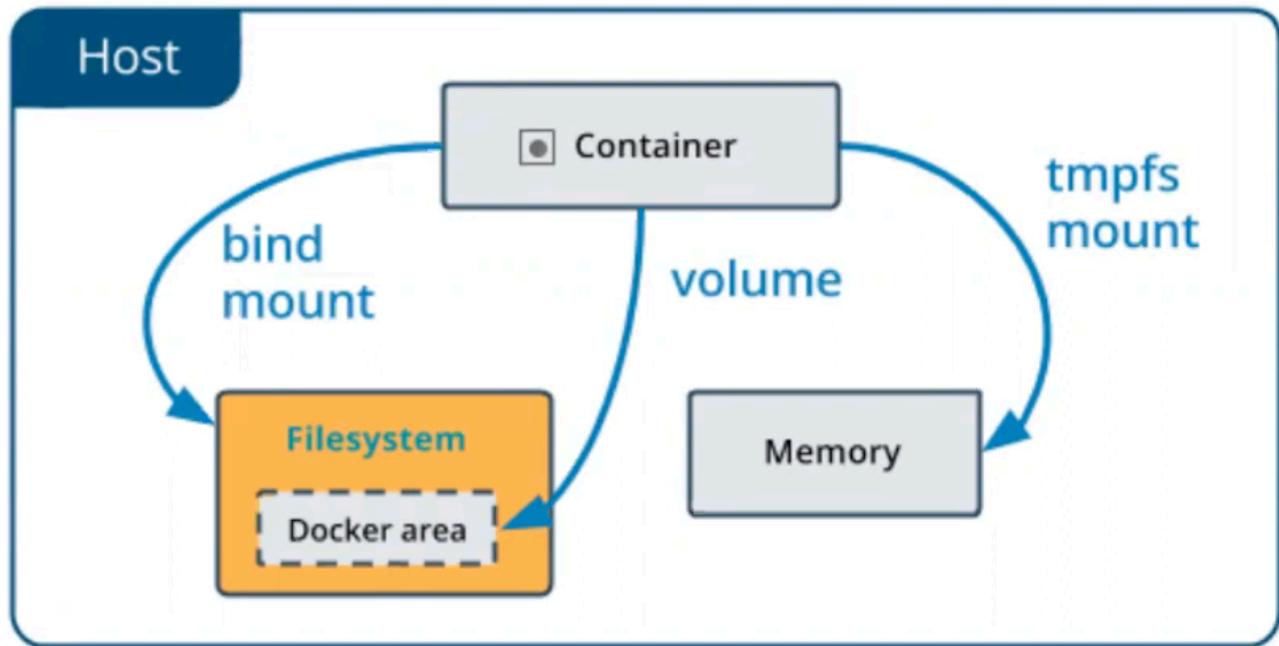
- **Download an image from a registry:** `docker image pull [OPTIONS] NAME[:TAG|@DIGEST]`.
- **Create and run a new container from image:** `docker run [OPTIONS] IMAGE [COMMAND] [ARG...]`.
 - `--name` : Assign name to container
 - `-it` : Combines `--interactive (-i)` and `--tty(-t)` : Keep STDIN open even if not attached, and allocate pseudo-TTY (interface)
 - `-v` : Bind mount a volume
 - `--rm` : Auto remove container when it exists
 - `--net/--network` : Connects a container to a network
 - `-p` : Publish a container's port(s) to the host
- **Build an image from a Dockerfile. A build's context is the set of files in the specified PATH or URL:** `docker build [OPTIONS] PATH | URL | -`.
 - `-f` : Name of the Dockerfile, default is 'Path/Dockerfile'
 - `--force-rm` : Always remove intermediate containers
 - `--no-cache` : Do not use cache when building the image.
- **List containers:** `docker ps [OPTIONS]`.
 - `-a` : Show all containers, running and stopped.
 - `-s` : Display total file size.
- **List images:** `docker images [OPTIONS] [REPOSITORY[:TAG]]`.
 - `-a` : Show all images, default hides intermediate images.
- **Remove one or more containers:** `docker rm [OPTIONS] CONTAINER [CONTAINER...]`.
 - `-f` : Force the removal of a running container (uses SIGKILL)
- **Remove one or more images:** `docker rmi [OPTIONS] IMAGE [IMAGE...]`.
 - `-f` : Force removal of the image.
- **Execute a command in a running container (usually to open an interactive bash in a running container):** `docker exec [OPTIONS] CONTAINER COMMAND [ARG...]`.
 - `-it` : Combines `--interactive (-i)` and `--tty(-t)` : Keep STDIN open even if not attached, and allocate pseudo-TTY
- **Copy folders/files between a container and the local filesystem (behaves like UNIX cp , having similar options for copying recursively, etc):**
 - `docker cp [OPTIONS] CONTAINER:SRC_PATH DEST_PATH` or
 - `docker cp [OPTIONS] SRC_PATH CONTAINER:DEST_PATH`.

Dockerfile

- **Dockerfile:** Dockerfiles are instructions for Docker to build images automatically. Using `docker build` users can create an automated build that executes several command-line instructions in succession (**in order from top to bottom**).
- General Syntax Format:
 - `# Comment`
 - `INSTRUCTION arguments`
- **FROM :** Create a new build stage from a base image:
 - Dockerfile must begin with a `FROM` instruction (may only be preceded by `ARG` instructions, which declare arguments that are used in `FROM` lines).
 - `FROM` instruction specifies the Parent Image from which you are building. E.g., `FROM ubuntu:24.04` .
- **RUN :** Executes build commands. Has two forms:
 1. Shell form: `RUN <command>` : command is run in shell.
 - Default shell is `/bin/sh -c` on linux, or `cmd /S /C` on windows.
 2. Exec form: `RUN ["executable", "param1", "param2"]` .
 - `git` dependencies, the directory you go into after `cd` when called by `RUN` **do not persist**. Only modifications to the file system will persist.
- **CMD :** Sets the command to be executed when running a container from an image. Has three forms:
 1. `CMD ["executable","param1","param2"]` (exec form, preferred)
 2. `CMD ["param1","param2"]` (exec form, as default parameters to `ENTRYPOINT`)
 3. `CMD command param1 param2` (shell form)
 - There can only be one `CMD` instruction in a Dockerfile. If you list more than one `CMD` , only the last one takes effect.
 - The purpose of a `CMD` is to provide defaults for an executing container. These defaults can include an executable, or they can omit the executable, in which case you must specify an `ENTRYPOINT` instruction as well.
- **ENV <key>=<value> [<key>=<value>...]** : Set environment variables.
 - Then `ENV` instruction sets the environment variable `<key>` to the value `<value>` . This value will be interpreted for other environment variables.
 - The environment variables set using `ENV` **will persist** when a container is run from the resulting image.
- **ARG <key>=<value> [<key>=<value>...]** : Use build-time variables.
 - Unlike `ENV` , `ARG` **will not persist** when a container is run from the resulting image.
- **COPY :** Copies new files or directories from `<src>` and adds them to the filesystem of the image at the path `<dest>` . Has two forms:

1. COPY [--chown=<user>:<group>] <src>... <dest>
 2. COPY [--chown=<user>:<group>] [<src>, ... <dest>]
 - [--chown=<user>:<group>] is an [OPTION] that is used to change the user and group ownership of files or directories.
- ENTRYPPOINT : Allows you to configure a container that will run as an executable. Has 2 forms:
 1. ENTRYPPOINT ["executable", "param1", "param2"] (exec form, preferred)
 2. ENTRYPPOINT command param1 param2 (shell form)
 - WORKDIR /path/to/workdir : Specifies the working directory.
 - Becomes the directory you'll be in when starting an interactive bash session.
 - If done before any of RUN , CMD , ENTRYPPOINT , COPY and ADD instructions, the work directory for them will be the path you give instead of / by default.
 - To learn more about these commands: <https://docs.docker.com/reference/dockerfile/>

Bind Mounts and Volume



- **Bind Mount:** Mounts an **existing file or directory** on the host machine into a container. The file or directory is referenced by its absolute path on the host machine.
- **Volume:** Creates a **new directory within Docker's storage directory** on the host machine, and Docker manages that directory's content.
- **Tmpfs (Temporary file system) mount:** Shares memory between container and host system.
- **Starting a container with a bind mount.** Two ways:
 1. -v : Combines all the options together in one field, separated by : .
 - First field is the path to the file/directory on the **host machine**.
 - Second field is the path where the file/directory is mounted in the **container**.

- Third field is optional, and is comma-separated list of options.
2. `--mount` : Consists of multiple key-value pairs, separated by commas each consisting of a `<key>=<value>` tuple. Some keys include:
- `type` , which can be `bind` , `volume` , or `tmpfs` .
 - `source` , the path to the file/directory on the **host**.
 - `destination` , the path to the file/directory mounted in the **container**.

Docker Network

- Container networking refers to the ability for containers to connect to and communicate with each other, and with non-Docker network services.
- Users can create a docker network to connect containers to them, or connect containers to non-Docker workloads.
- Create a user-defined bridge network: `$ docker network create <name>`
- Connect a container to a network: `$ docker network connect <net-name> <container>`
- Use the `--network host` option when using `docker run` to share the host's network with container.
- To learn more: <https://www.youtube.com/watch?v=zJD7QYQtIKc>

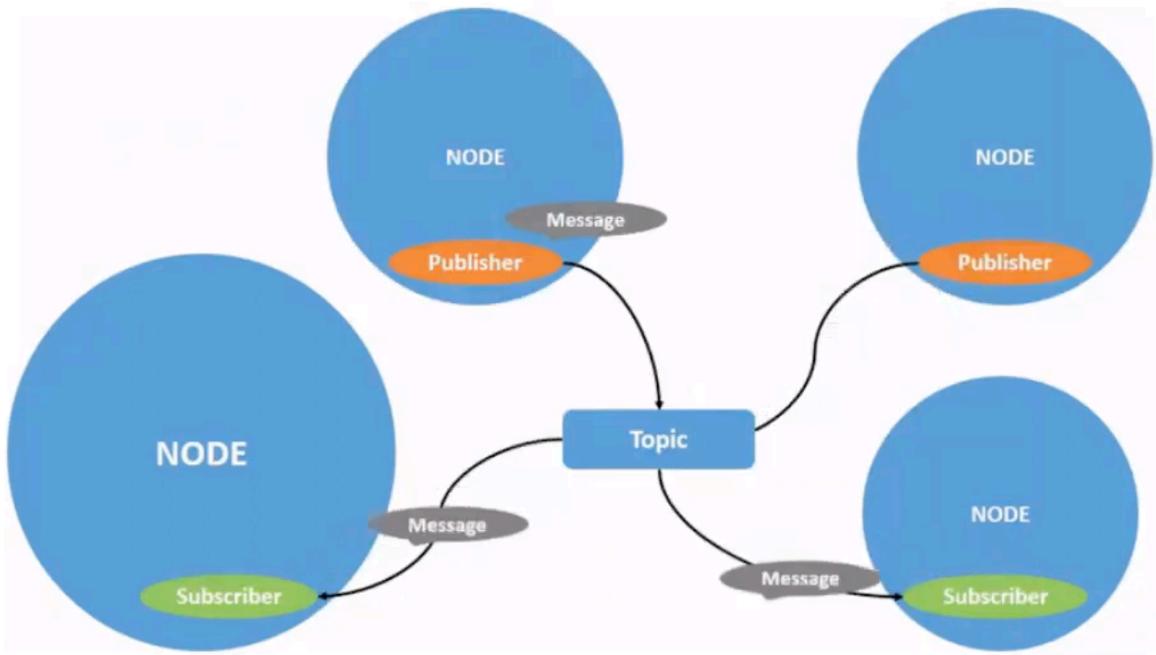
Docker Compose

- Compose is a tool for defining and running multiple containers at the same time.
- A yaml file `docker-compose.yml` is used to configure.
- Run `docker compose up` or `docker-compose up` (if you've installed the docker-compose) to start all your containers.

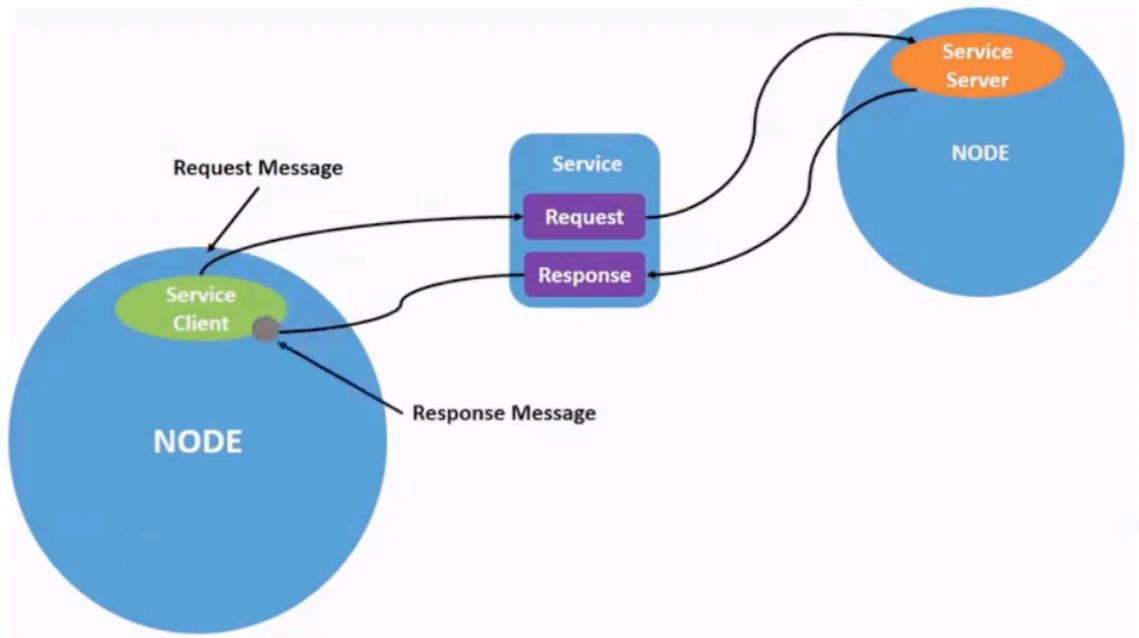
ROS 2

- **Robot Operating System (ROS) 2:**
 - **Distributed:** Utilizes a graph-like structure where nodes, which are individual software processes, communicate with each other through topics, services, and actions.
 - **Peer to peer:** Nodes communicate directly with each other
 - **Multi-lingual:** Supports C++ and Python
 - **Light-weight**
 - **Free and open-source**
- **ROS Graph:** The ROS graph is a network of ROS 2 elements processing data together at one time. It encompasses all executables and the connections between them if you were to map them all out and visualize them.
- **Nodes:**

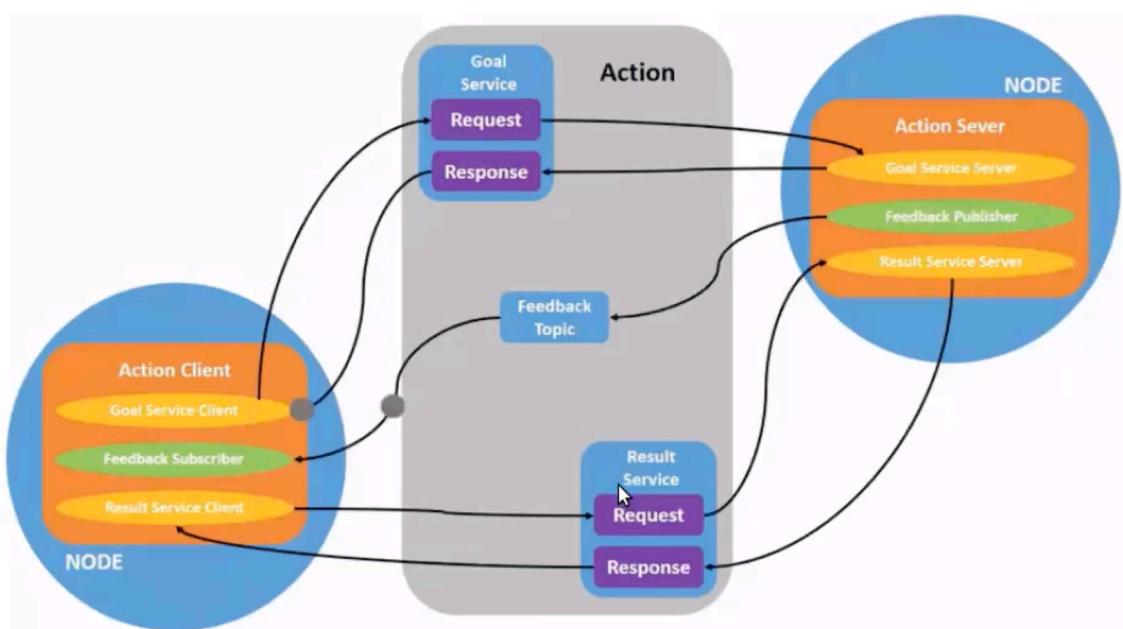
- Each node in ROS 2 should be responsible for a single, module purpose (e.g., one node for controlling wheel motors, one node for controlling a laser range-finder, etc).
- Each node can send and receive data to other nodes via topics, services, actions, or parameters.
- To run a node: `ros 2 run <package_name> <executable_name>`
- To list the nodes that are currently running: `ros2 node list`
- To get information about a given node: `node info <node_name>`
- **Topics:** How nodes communicate with each other (main way data is moved between nodes). Acts as a bus for nodes to exchange messages.
 - A node **publishes** (sends) data to **topics**.
 - A node **subscribes** (receives) data from **topics**.
 - Topics don't have to only be point-to-point communication; it can be one-to-many, many-to-one, or many-to-many.
 - **Graph of how nodes and topics are communicating with each other, also shows the types of messages:** `rqt_graph`
 - **Lists all active topics:** `ros2 topic list`
 - **Lists all active topics with their message types:** `ros2 topic list -t`
 - **Displays messages being published on a topic:** `ros2 topic echo <topic_name>`
 - **Shows detailed information about a topic:** `ros2 topic info <topic_name>`
 - **Displays the structure of a message type:** `ros2 interface show <msg_type>`
 - **Publishes a message to a topic from the command line:** `ros2 topic pub <topic_name> <msg_type> '<args>'`
 - **Displays the publishing frequency of a topic:** `ros2 topic hz <topic_name>`
 - Diagram:



- **Services:** Services are another method of communication for nodes in the ROS graph. Services are based on call-and-response model, versus topics' publisher-subscriber model. While topics allow nodes to subscribe to data streams and get continual updates, **services only provide data when they are specifically called by a client.**
 - Can be many service clients using the same service. But **only one** service server for a service.
 - **Lists all active services:** `ros2 service list`
 - **Shows the type of a specific service:** `ros2 service type <service_name>`
 - **Lists all active services with their types:** `ros2 service list -t`
 - **Finds services by service type:** `ros2 service find <type_name>`
 - **Displays the structure of a service type (what type of data it needs for a request and what for response):** `ros2 interface show <type_name>.srv`
 - **Calls a service from the command line:** `ros2 service call <service_name> <service_type> '<args>'`
 - Diagram:



- **Actions:** Actions are type of communication method meant for **long running tasks**. Consist of: a goal, feedback, and a result.
 - Actions are built on topics and services.
 - Unlike services, actions are preemptable (you can cancel them while executing) and provide steady feedback.
 - An "action client" node sends a goal to an "action server" node that acknowledges the goal and returns a stream of feedback and a result.
 - Diagram:



- **Parameters:** Configuration value for nodes (node settings).

- Each node maintains its own parameters (such as `int`, `float`, `bool`, etc).
- All parameters are dynamically reconfigurable and built off ROS 2 services.
- **Lists all parameters:** `ros2 param list`
- **Gets the value of a parameter:** `ros2 param get <node_name> <parameter_name>`
- **Sets the value of a parameter:** `ros2 param set <node_name> <parameter_name> <value>`
- Can also use YAML or launch file to define parameters: Useful to have all parameters in one place while tuning & set parameters for multiple nodes.
- Set in C++: `this->declare_parameter("my_parameter", "world");`
- Get in C++: `this->get_parameter("my_parameter");`
- Learn more here: <https://roboticsbackend.com/rclcpp-params-tutorial-get-set-ros2-params-with-cpp/> & <https://docs.ros.org/en/foxy/Tutorials/Beginner-Client-Libraries/Using-Parameters-In-A-Class-CPP.html>
- **Workspace:** Directory containing all ROS 2 packages.
 - Before using ROS 2, you must source your ROS 2 installation workspace in the terminal you plan to work in to make them available for you to use in that terminal.
 - **Underlay:** Main ROS 2 environment.
 - Sourcing an **Overlay**: A secondary workspace where you can add new packages without interfering with the existing ROS 2 workspace that you are extending, or "underlay".
 - Your underlay must contain all the dependencies of all the packages in your overlay.
 - Packages in overlay override packages in underlay.
 - Possible to have several layers of underlays and overlays - where each successive overlay uses the packages of its parent underlays.
 - **Create a new workspace:** `$ mkdir -p <your_workspace>/src`. Then put desired ROS 2 packages inside `src`
 - **Resolving dependencies:** `$ cd <your_workspace> $ rosdep install -i --from-path src --rosdistro foxy -y`. This will install dependencies declared in `package.xml` from all packages in the `src` directory for ROS 2 foxy.
 - **Build workspace with colcon** : From the root of your workspace: `$ colcon build`.
 - `--packages-up-to` : Build packages you want, plus all its dependencies, but not the whole workspace. Saves time if you don't need all packages in workspace
 - Once build finishes, you should see: `build`, `install`, `log`, `src` directories in your workspace. The `install` directory is where your workspace's setup files are.
 - **Sourcing an Underlay**: Give your underlay access to your packages on top of the base ROS 2 environment.
 - Do it by writing in new terminal: `$ source /opt/ros/foxy/setup.bash`

- Then, in root of desired workspace, `$ cd <your_workspace>` and `$ source isntall/local_setup.bash`
- Instead of doing the above two notes you could combine them with: `$ install/setup.bash` in your workspace.
- **ROS 2 Packages:** A package can be considered a container in your ROS 2 code. To install and share your code, you'll need to organize it in a package.
 - **Package Creation:** ROS 2 uses **ament** as its build system and **colcon** as its build tool. Use CMake for your package.
 - CMake package requires:
 - File containing meta info about the package: `package.xml`
 - File describing how to build the code within the package: `CMakeLists.txt`
 - Most simple structure looks like folder named `my_package` with the above two files inside.
 - A single workspace can contain **many packages**, each in their own folder. You **cannot have nested packages**.
 - Best practice: Have `src` folder within your workspace and create all your packages in there.
 - **Create Package:** `$ cd <your_workspace>/src` and `$ ros 2 pkg create --build-type ament_cmake <package_name>`
 - **Package Contents:**
 - Contains files: `CMakeLists.txt`, and `package.xml`, and folders: `include`, `src`. Node source files (`.cpp`) are in `src`, and header files (`.hpp`) are in `include`.
 - **Customizing package.xml:** Fill in name and email on `maintainer` line, edit description to summarize the package, update the `license` line. Fill in your dependencies under the `_depend` tags. For more docs about what types of depend tags: <https://docs.ros.org/en/humble/Tutorials/Intermediate/Rosdep.html>
- Wow you made it all the way down here! If you need to learn more about anything covered above, you can always refer to the docs. Our club has a repository where we will continually add documentation and tutorials for you to use here: <https://github.com/sfu-racerbot/racerbot-docs/tree/main>